

# Building Autonomous AI Trading Bots with Claude Code CLI

Claude Code provides a powerful, extensible platform for building autonomous trading systems through **hooks** (deterministic event-driven automation), **custom tools** (MCP server integration), **persona systems** (CLAUDE.md configuration), and **external orchestration** (multi-agent coordination). This report details the complete technical architecture needed to integrate Claude Code with QuantConnect LEAN for autonomous options trading development.

## Core extensibility: hooks, tools, and MCP servers

Claude Code's extensibility rests on three pillars that work together to create sophisticated autonomous development workflows. **Hooks** provide deterministic control over Claude's actions at 8 lifecycle points, while **MCP servers** extend Claude's capabilities with custom tools, and **SDK tools** enable in-process extensions.

### Hook architecture and event types

Hooks are shell commands configured in `(.claude/settings.json)` that execute at specific lifecycle points. `(Claude)` The configuration hierarchy follows priority order: enterprise managed (`(/etc/claude-code/managed-settings.json)`), user settings (`(~/.claude/settings.json)`), project settings (`(.claude/settings.json)`), and local project (`(.claude/settings.local.json)`).

| Hook Event                       | Fires When                 | Blocking       | Trading Use Case                                     |
|----------------------------------|----------------------------|----------------|--|
| <code>(PreToolUse)</code>        | Before tool execution      | Exit 2 or JSON | Block unauthorized trades, validate order parameters |
| <code>(PostToolUse)</code>       | After tool completion      | No             | Log executions, trigger alerts, validate results     |
| <code>(UserPromptSubmit)</code>  | User submits prompt        | Exit 2         | Inject market context, validate trading instructions |
| <code>(PermissionRequest)</code> | Permission dialog shown    | JSON           | Auto-approve/deny specific operations                |
| <code>(Stop)</code>              | Claude finishes responding | JSON           | Ensure order confirmations, force continuation       |
| <code>(SessionStart)</code>      | Session starts/resumes     | No             | Load positions, market context                       |

The hook input format passes comprehensive context via stdin JSON including `(Claude)` `(session_id)`, `(transcript_path)`, `(hook_event_name)`, `(tool_name)`, and `(tool_input)`. `(Claude)` Exit codes control behavior: **0** allows execution (stdout shown), **2** blocks with feedback to Claude (stderr passed automatically), and other codes generate non-blocking errors.

A trading-specific risk validator hook demonstrates the pattern:

```
json
```

```
{
  "hooks": {
    "PreToolUse": [
      {
        "matcher": "mcp_broker_execute*mcp_broker_place*",
        "hooks": [
          {
            "type": "command",
            "command": "python3 \"$CLAUDE_PROJECT_DIR/.claude/hooks/risk_validator.py\"",
            "timeout": 10
          }
        ]
      },
      {
        "PostToolUse": [
          {
            "matcher": "mcp_broker_*",
            "hooks": [
              {
                "type": "command",
                "command": "python3 \"$CLAUDE_PROJECT_DIR/.claude/hooks/log_trade.py\""
              }
            ]
          }
        ]
      }
    }
  }
}
```

## MCP server integration for market data and execution

MCP (Model Context Protocol) servers extend Claude with custom tools ([anthropic](#)) for market data fetching, order execution, and backtesting. Several trading-relevant MCP servers exist: **Alpha Vantage** provides stocks, options, forex, and technical indicators; **Polygon.io** offers 35+ tools covering stocks, options, forex, and crypto with full OPRA feed access; **Twelve Data** delivers real-time quotes and WebSocket streaming.

Configure MCP servers via CLI or JSON:

```
bash

# Add market data server
claude mcp add --transport http alphavantage https://mcp.alphavantage.co/mcp?apikey=KEY

# Add custom trading server with environment variables
claude mcp add broker-api --env BROKER_API_KEY=xxx -- python3 ./trading_mcp_server.py
```

For team sharing, create `.mcp.json` in the project root: [Anthropic](#)

```
json
```

```
{
  "mcpServers": {
    "financial-datasets": {
      "command": "python3",
      "args": ["./mcp/trading_server.py"],
      "env": {
        "BROKER_API_KEY": "${BROKER_API_KEY}",
        "BROKER_ACCOUNT_ID": "${BROKER_ACCOUNT_ID}"
      }
    }
  }
}
```

Custom trading MCP servers can implement tools for options chain retrieval, Greek calculations, spread order placement, and IV surface analysis. Tools follow the naming convention `mcp_{server_name}_{tool_name}`, allowing precise hook matching for risk controls.

## Persona system and CLAUDE.md configuration

CLAUDE.md serves as the "constitution" for Claude Code behavior, loaded automatically at every session start. The hierarchical system supports global (`~/claude/CLAUDE.md`), project (checked into git), local (gitignored for personal overrides), and nested folder-specific configurations.

### Implementing role-based personas

Multiple methods enable persona switching for different development tasks:

**Command-line system prompt injection** provides the most direct approach:

```
bash

# Senior coder persona
claude --append-system-prompt "You are a senior software engineer with 15 years experience. Focus on clean architecture, SC

# Code reviewer persona
claude --append-system-prompt "You are conducting senior-level code review. Focus on: security vulnerabilities, performance

# QA specialist persona
claude --append-system-prompt "You are a QA automation specialist. Focus on comprehensive test coverage, edge cases, bou
```

**Custom subagents** (native Claude Code feature) provide isolated context windows and specialized expertise.

Create files in `.claude/agents/`:

```
markdown
```

```
---
```

```
name: risk-reviewer
description: Risk and compliance reviewer for trading systems
tools: Read, Grep, Glob
model: opus
```

```
---
```

You are a risk management specialist reviewing trading code.

### ## Review Checklist

- Position size limits enforced
- Stop-loss mechanisms implemented
- Kill switch functionality present
- Circuit breakers configured
- Audit logging complete

**Shell function personas** combine CLI flags with MCP configurations: [Decoding](#)

```
bash
```

```
# Research persona with specialized tools
ccr() {
    claude ${*} --dangerously-skip-permissions \
    --model opus \
    --mcp-config ~/.claude/mcp-research.json \
    --append-system-prompt "You are a code research and analysis specialist..."
}
```

## CLAUDE.md structure for trading projects

```
markdown
```

## # Trading Bot Development Guidelines

### ## STRICT REQUIREMENTS

- NEVER execute live trades without explicit confirmation
- ALL trading operations require paper trading validation first
- Maximum position size: 2% of portfolio per trade
- Mandatory stop-loss on all positions

### ## Architecture

- Microservices with event-driven communication
- FastAPI for API services
- PostgreSQL for persistence, Redis for caching

### ## Testing Requirements

- Minimum 80% coverage
- Backtest validation for all trading strategies
- Integration tests for API endpoints

### ## Safety

- Never commit API keys
- Implement rate limiting on public endpoints
- Log all trade decisions with reasoning

Use `@path/to/import` syntax to modularly import documentation, patterns, and specifications into CLAUDE.md files. [Claude](#) [Shipyard](#)

## External scaffolding and orchestration systems

External scaffolding enables multi-hour autonomous sessions, parallel agent coordination, and persistent workflows essential for trading system development.

### Claude-Flow: enterprise multi-agent orchestration

**Claude-Flow** provides queen-led AI coordination with specialized worker agents, 100+ MCP tools, and AgentDB integration with **96x-164x faster vector search**:

```
bash
```

```

# Initialize orchestration
npx claude-flow@alpha init --force
claude mcp add claude-flow npx claude-flow@alpha mcp start

# Multi-agent trading bot development
npx claude-flow@alpha hive-mind spawn "market-data-service" --namespace market --claude
npx claude-flow@alpha hive-mind spawn "order-execution" --namespace orders --claude
npx claude-flow@alpha hive-mind spawn "risk-management" --namespace risk --claude

```

**CLI Agent Orchestrator (CAO)** from AWS Labs provides hierarchical supervision with tmux session isolation, supporting handoff (synchronous), assign (asynchronous), and message-passing patterns. (AWS)

## Docker containerization patterns

Anthropic's official devcontainer provides secure isolated execution with firewall configuration allowing only npm registry, GitHub, and Claude API access while denying all other external connections.

**Claude Code Sandbox** enables safe `--dangerously-skip-permissions` execution:

```

json

{
  "dockerImage": "claude-code-sandbox:latest",
  "detached": false,
  "autoPush": true,
  "autoCreatePR": true,
  "setupCommands": ["npm install", "npm run build"],
  "mounts": [
    {
      "source": "./data",
      "target": "/workspace/data",
      "readonly": false
    }
  ],
  "allowedTools": ["*"],
  "maxThinkingTokens": 100000,
  "bashTimeout": 600000
}

```

## Extended workflow patterns for overnight automation

**Continuous Claude** implements automated PR loops—creating branches, running Claude Code, committing changes, pushing PRs, waiting for CI checks, and merging:

```
bash
```

```
continuous-claude --prompt "implement backtesting engine" \
--max-runs 10 \
--max-cost 25.00 \
--owner YourOrg \
--repo trading-bot \
--merge-strategy squash
```

Context persists across iterations via `(SHARED_TASK_NOTES.md)`, enabling multi-session progress tracking.

**Cron-based scheduling** enables overnight autonomous development:

```
yaml

# Kubernetes CronJob for daily strategy analysis
apiVersion: batch/v1
kind: CronJob
metadata:
  name: claude-autonomous-scheduler
spec:
  schedule: "0 6 * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: claude-runner
              image: claude-code:latest
              command: ["claude", "--auto-approve", "--batch", "Analyze trading strategies"]
```

**Git worktrees** enable parallel development across multiple features simultaneously. [anthropic](#)

## Compliance, safety, and forcing mechanisms

Trading systems require deterministic guardrails that cannot be bypassed by AI decision-making.

### Permission configuration and guardrails

```
json
```

```
{  
  "permissions": {  
    "allow": [  
      "mcp__market-data__*",  
      "Bash(python3 analyze_.py)"  
    ],  
    "ask": [  
      "mcp__broker__execute*",  
      "mcp__broker__place*"  
    ],  
    "deny": [  
      "Read(./.env*)",  
      "Bash(rm -rf:*)",  
      "WebFetch"  
    ]  
  }  
}
```

**Sandboxing** (enabled via `(/sandbox)`) provides OS-level isolation using Linux bubblewrap or macOS seatbelt, reducing permission prompts by **84%** while preventing prompt injection attacks from stealing credentials or exfiltrating data. [Anthropic](#) [anthropic](#)

## Trading-specific safety mechanisms

**Kill switches** implement multi-level emergency halt:

```
python
```

```

class TradingKillSwitch:
    def __init__(self, config):
        self.max_drawdown_percent = config.get('max_drawdown', 5.0)
        self.max_daily_loss = config.get('max_daily_loss', 5000)

    def check_and_halt(self, portfolio_state):
        if portfolio_state.drawdown_percent > self.max_drawdown_percent:
            self.halt_trading("Max drawdown exceeded")
            return True
        if portfolio_state.daily_pnl < -self.max_daily_loss:
            self.halt_trading("Max daily loss exceeded")
            return True
        return False

    def halt_trading(self, reason):
        cancel_all_open_orders()
        send_alert(f"KILL SWITCH ACTIVATED: {reason}")

```

**Circuit breakers** implement failure threshold patterns with CLOSED, OPEN, and HALF\_OPEN states, auto-resetting after configured timeout periods.

**Position limits** enforce maximum position size (recommend **2%** per position), maximum sector exposure (**20%**), and total portfolio exposure limits before every order.

**Paper trading validation** requirements before live deployment:

- 100+ backtested trades
- 2-4 weeks minimum paper trading [Tradetron](#)
- Consistent positive expectancy demonstrated
- Win rate, drawdown, and Sharpe ratio documented
- Emergency procedures tested

## Checkpointing and recovery

Claude Code's built-in checkpointing automatically tracks file edits, enabling [/rewind](#) to restore conversation state, code state, or both. [ClaudeLog](#) Press [Esc](#) twice for quick rewind access.

Third-party tools like **Claude Code Rewind** provide enhanced session timeline visualization, diff viewing, and rollback with preview capabilities. [GitHub](#)

## QuantConnect LEAN integration

LEAN is QuantConnect's professional-caliber algorithmic trading engine [github](#) supporting Python 3.11 and

C#.

## Local development setup

```
bash

pip install lean
lean login # Requires QuantConnect membership
lean init # Creates lean.json and data/ directory
```

## Algorithm structure

```
python

from AlgorithmImports import *

class OptionsTradingAlgorithm(QCAlgortithm):
    def Initialize(self):
        self.SetStartDate(2022, 1, 1)
        self.SetEndDate(2022, 12, 31)
        self.SetCash(100000)

        # Add underlying and options
        self.symbol = self.AddEquity("SPY", Resolution.Minute).Symbol
        option = self.AddOption("SPY")
        option.SetFilter(minStrike=2, maxStrike=2, minExpiry=0, maxExpiry=30)

    def OnData(self, slice):
        # Trading logic implementation
        pass
```

## Development workflow

```
bash
```

```

# Local backtest
lean backtest "MyStrategy" --output ./backtests/

# Cloud validation
lean cloud backtest "MyStrategy" --push --open

# Paper trading deployment
lean live deploy "MyStrategy" --brokerage "Paper Trading"

# Live deployment to Schwab
lean live deploy "MyStrategy" --brokerage "Charles Schwab" \
--schwab-app-key "KEY" --schwab-secret "SECRET"

```

QuantConnect supports **Charles Schwab**, Interactive Brokers, TradeStation, Alpaca, Tradier, Tastytrade, and multiple crypto exchanges for live deployment. (QuantConnect)

## Charles Schwab API integration

The Schwab API requires registration at developer.schwab.com (separate from brokerage), app creation with callback URL configuration, and OAuth 2.0 authentication.

Using the unofficial schwab-py wrapper:

```

python

from schwab import auth, client

api_key = 'YOUR_API_KEY'
app_secret = 'YOUR_APP_SECRET'
callback_url = 'https://127.0.0.1:8182/'
token_path = '/path/to/token.json'

# Authentication opens browser for Schwab login
c = auth.easy_client(api_key, app_secret, callback_url, token_path)

# Get historical data
r = c.get_price_history_every_day('AAPL')

```

**Key limitations:** Access tokens expire after **30 minutes**, refresh tokens after **7 days** requiring manual re-authentication, and order rate limits of **120 orders/minute**.

## CI/CD integration patterns

The official Claude Code GitHub Action enables automated code review, implementation, and maintenance:

```

yaml

name: Claude PR Review
on:
  pull_request:
    types: [opened, synchronize]

jobs:
  review:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: anthropics/clause-code-action@v1
        with:
          anthropic_api_key: ${{ secrets.ANTHROPIC_API_KEY }}
          prompt: |
            Review this PR for code quality, security vulnerabilities,
            performance issues, and test coverage.

```

**Headless mode** supports CI/CD automation: [anthropic](#)

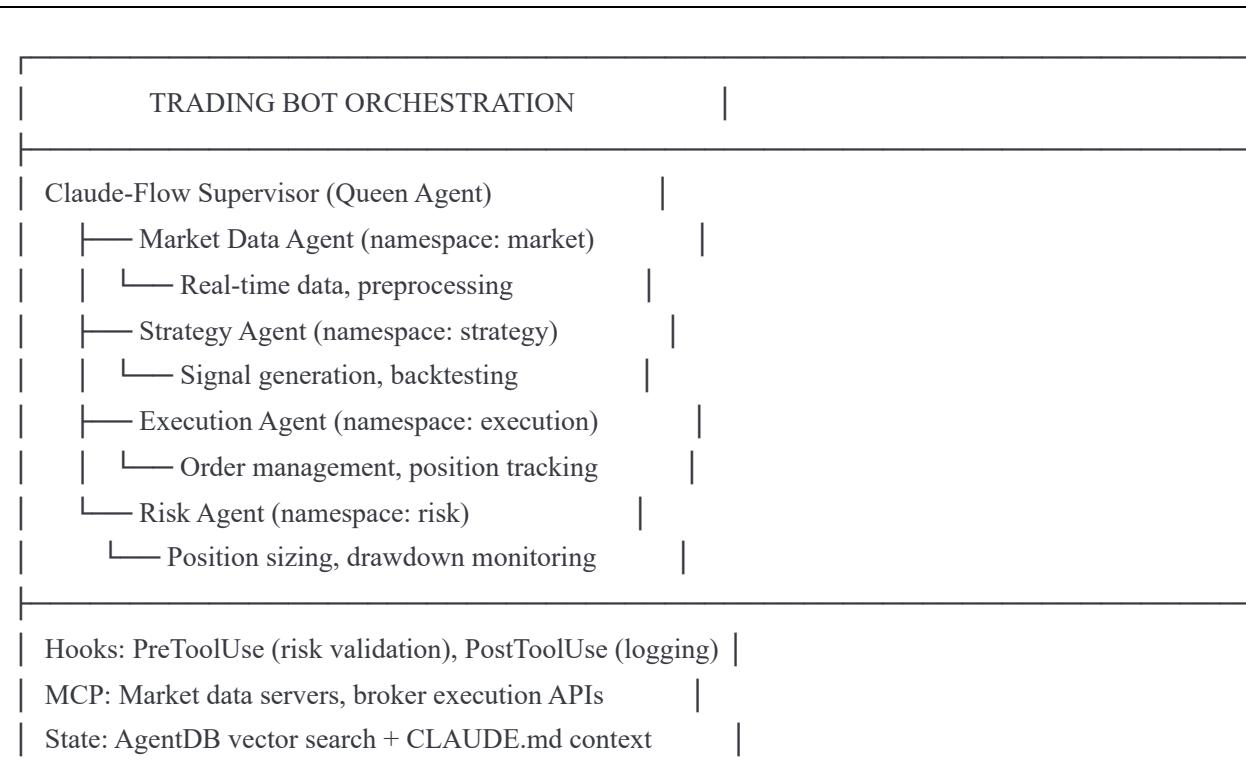
```

bash

claude -p "Generate tests for auth module" --output-format stream-json | jq '.type'

```

## Recommended architecture for autonomous trading development



| CI/CD: GitHub Actions + Claude Code Action  
| Runtime: Docker containers with network isolation

This architecture provides **deterministic risk controls** via hooks (cannot be bypassed by AI), **real-time market data** via MCP servers, **order execution** with multi-layer validation, **audit logging** via PostToolUse hooks, and **session context** via SessionStart hooks loading positions and market state.

## Conclusion

Building autonomous trading bots with Claude Code requires layering multiple systems: **hooks** for deterministic safety controls, **MCP servers** for market data and execution capabilities, **CLAUDE.md** for project context and constraints, and **external orchestration** (Claude-Flow or CAO) for multi-agent coordination. The key insight is that trading safety mechanisms must exist in hooks—deterministic code that executes regardless of AI decisions—rather than relying on prompt engineering alone.

For production deployment, implement the **complete safety stack**: PreToolUse hooks for order validation, position limits enforced at the code level, circuit breakers with automatic halt capabilities, kill switches accessible both programmatically and manually, and comprehensive audit logging of all trading decisions with reasoning. Paper trade for a minimum of 2-4 weeks ([Tradetron](#)) with 100+ simulated trades before any live deployment, and always run Claude Code in sandboxed Docker containers with network isolation for development work.